

Different_MLP_architectures_on_MNIST_dataset

January 21, 2019

```
In [0]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use t
from keras.utils import np_utils
import pandas as pd
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
import keras
keras.backend.backend()
```

Using TensorFlow backend.

```
Out[0]: 'tensorflow'
```

```
In [0]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
In [0]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 1s 0us/step

```
In [0]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%)")
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%)")
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 3 dimensional vector
        # for each image we have a (28*28) vector
        # we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [0]: # after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%)")
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%)")
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

```
In [0]: # An example data point
        print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150 27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
```

```

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 249 253 249 64 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 46 130 183 253
253 207 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 39 148 229 253 253 253 250 182 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 24 114 221 253 253 253
253 201 78 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 23 66 213 253 253 253 253 198 81 2 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 18 171 219 253 253 253 253 195
80 9 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
55 172 226 253 253 253 253 244 133 11 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 136 253 253 253 212 135 132 16
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0]

```

In [0]:

```

In [0]: # if we observe the above matrix each cell is having a value between 0-255
        # before we move to apply machine learning algorithms lets try to normalize the data
        #  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

```

```

X_train = X_train/255
X_test = X_test/255

```

```

In [0]: # example data point after normlizing
        print(X_train[0])

```

```

[0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.]

```

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.01176471	0.07058824	0.07058824	0.07058824
0.49411765	0.53333333	0.68627451	0.10196078	0.65098039	1.
0.96862745	0.49803922	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686
0.99215686	0.95686275	0.52156863	0.04313725	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.53333333	0.99215686

[illegible]

```
In [0]: # here we are having a class number for each image
        print("Class label of first image :", y_train[0])

        # lets convert this into a 10 dimensional vector
        # ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
        # this conversion needed for MLPs

        Y_train = np_utils.to_categorical(y_train, 10)
        Y_test = np_utils.to_categorical(y_test, 10)

        print("After converting the output into a vector : ", Y_train[0])

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Softmax classifier

```
In [0]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor.

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])
```

```

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

```
In [0]: # some model parameters
```

```

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128

```

```
nb_epoch = 20
```

```
In [0]: # start building a model
```

```
model = Sequential()
```

```
# The model needs to know what input shape it should expect.  
# For this reason, the first layer in a Sequential model  
# (and only the first, because following layers can do automatic shape inference)  
# needs to receive information about its input shape.  
# you can use input_shape and input_dim to pass the shape of input
```

```
# output_dim represent the number of nodes need in that layer  
# here we have 10 nodes
```

```
model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

```
In [0]: # Before training a model, you need to configure the learning process, which is done v
```

```
# It receives three arguments:  
# An optimizer. This could be the string identifier of an existing optimizer , https://  
# A loss function. This is the objective that the model will try to minimize., https://  
# A list of metrics. For any classification problem you will want to set this to metri
```

```
# Note: when using the categorical_crossentropy loss, your targets should be in categor  
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional v  
# for a 1 at the index corresponding to the class of the sample).
```

```
# that is why we converted out labels into vectors
```

```
model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
# Keras models are trained on Numpy arrays of input data and labels.  
# For training a model, you will typically use the fit function
```

```
# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, vali  
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_e  
# validation_steps=None)
```

```
# fit() function Trains the model for a fixed number of epochs (iterations on a datase
```

```
# it returns A History object. Its History.history attribute is a record of training l  
# metrics values at successive epochs, as well as validation loss values and validation
```

```
# https://github.com/openai/baselines/issues/20
```

```
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=
```

Train on 60000 samples, validate on 10000 samples


```

Epoch 1/20
60000/60000 [=====] - 2s 27us/step - loss: 1.2892 - acc: 0.6962 - val.
Epoch 2/20
60000/60000 [=====] - 1s 19us/step - loss: 0.7151 - acc: 0.8415 - val.
Epoch 3/20
60000/60000 [=====] - 1s 19us/step - loss: 0.5861 - acc: 0.8599 - val.
Epoch 4/20
60000/60000 [=====] - 1s 20us/step - loss: 0.5245 - acc: 0.8690 - val.
Epoch 5/20
60000/60000 [=====] - 1s 21us/step - loss: 0.4870 - acc: 0.8756 - val.
Epoch 6/20
60000/60000 [=====] - 1s 20us/step - loss: 0.4612 - acc: 0.8800 - val.
Epoch 7/20
60000/60000 [=====] - 1s 21us/step - loss: 0.4422 - acc: 0.8836 - val.
Epoch 8/20
60000/60000 [=====] - 1s 19us/step - loss: 0.4273 - acc: 0.8871 - val.
Epoch 9/20
60000/60000 [=====] - 1s 19us/step - loss: 0.4153 - acc: 0.8892 - val.
Epoch 10/20
60000/60000 [=====] - 1s 17us/step - loss: 0.4053 - acc: 0.8910 - val.
Epoch 11/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3969 - acc: 0.8932 - val.
Epoch 12/20
60000/60000 [=====] - 1s 18us/step - loss: 0.3897 - acc: 0.8946 - val.
Epoch 13/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3833 - acc: 0.8960 - val.
Epoch 14/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3777 - acc: 0.8971 - val.
Epoch 15/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3727 - acc: 0.8981 - val.
Epoch 16/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3682 - acc: 0.8991 - val.
Epoch 17/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3641 - acc: 0.9000 - val.
Epoch 18/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3604 - acc: 0.9007 - val.
Epoch 19/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3569 - acc: 0.9017 - val.
Epoch 20/20
60000/60000 [=====] - 1s 17us/step - loss: 0.3538 - acc: 0.9021 - val.

```

```

In [0]: score_in_out = model.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_in_out[0])
        print('Test accuracy:', score_in_out[1])

        fig,ax = plt.subplots(1,1)
        ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

```

```

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of e

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.33589279960989954

Test accuracy: 0.9093



Observations 1. We did not use any hidden layer and have only input and output layer and got accuracy ~90% which is best.

1 1. MLP + ReLU + ADAM with 2 hidden layers

```
In [0]: model_relu = Sequential()
        model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0,
        model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
        model_relu.add(Dense(output_dim, activation='softmax'))

        print(model_relu.summary())

        model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

        history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

```
-----
Layer (type)                 Output Shape              Param #
=====
dense_2 (Dense)              (None, 512)              401920
-----
dense_3 (Dense)              (None, 128)              65664
-----
dense_4 (Dense)              (None, 10)               1290
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
-----
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 6s 108us/step - loss: 0.2267 - acc: 0.9321 - val_loss: 0.1044 - val_acc: 0.9751
Epoch 2/20
60000/60000 [=====] - 6s 104us/step - loss: 0.0849 - acc: 0.9751 - val_loss: 0.0444 - val_acc: 0.9884
Epoch 3/20
60000/60000 [=====] - 6s 106us/step - loss: 0.0540 - acc: 0.9832 - val_loss: 0.0374 - val_acc: 0.9884
Epoch 4/20
60000/60000 [=====] - 6s 102us/step - loss: 0.0374 - acc: 0.9884 - val_loss: 0.0255 - val_acc: 0.9921
Epoch 5/20
60000/60000 [=====] - 6s 102us/step - loss: 0.0255 - acc: 0.9921 - val_loss: 0.0209 - val_acc: 0.9931
Epoch 6/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0209 - acc: 0.9931 - val_loss: 0.0166 - val_acc: 0.9948
Epoch 7/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0166 - acc: 0.9948 - val_loss: 0.0140 - val_acc: 0.9951
Epoch 8/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0140 - acc: 0.9951 - val_loss: 0.0125 - val_acc: 0.9960
Epoch 9/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0125 - acc: 0.9960 - val_loss: 0.0106 - val_acc: 0.9970
Epoch 10/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0106 - acc: 0.9970 - val_loss: 0.0106 - val_acc: 0.9970
```

```

Epoch 11/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0142 - acc: 0.9952 - va
Epoch 12/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0092 - acc: 0.9970 - va
Epoch 13/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0132 - acc: 0.9960 - va
Epoch 14/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0103 - acc: 0.9963 - va
Epoch 15/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0090 - acc: 0.9972 - va
Epoch 16/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0055 - acc: 0.9982 - va
Epoch 17/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0090 - acc: 0.9971 - val
Epoch 18/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0088 - acc: 0.9972 - va
Epoch 19/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0082 - acc: 0.9975 - val
Epoch 20/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0057 - acc: 0.9983 - val

```

```

In [0]: score_relu2_adam = model_relu.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu2_adam[0])
        print('Test accuracy:', score_relu2_adam[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0938020015134396

Test accuracy: 0.9807



```
In [0]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
```

```
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
violin_data = remove_na(group_data)
```



2 MLP + ReLU + ADAM with 2 hidden layers

```
In [0]: model_relu = Sequential()
```

```
model_relu.add(Dense(672, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0,
```

```
model_relu.add(Dense(325, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
```

```
model_relu.add(Dense(output_dim, activation='softmax'))
```

```
print(model_relu.summary())
```

```
-----
Layer (type)                Output Shape          Param #
```

```

=====
dense_5 (Dense)                (None, 672)                527520
-----
dense_6 (Dense)                (None, 325)                218725
-----
dense_7 (Dense)                (None, 10)                 3260
=====
Total params: 749,505
Trainable params: 749,505
Non-trainable params: 0
-----
None

```

```
In [0]: model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)
```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 10s 165us/step - loss: 0.2132 - acc: 0.9354 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 2/20
60000/60000 [=====] - 9s 157us/step - loss: 0.0761 - acc: 0.9764 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 3/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0465 - acc: 0.9849 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 4/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0335 - acc: 0.9894 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 5/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0227 - acc: 0.9928 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 6/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0200 - acc: 0.9929 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 7/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0183 - acc: 0.9935 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 8/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0181 - acc: 0.9941 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 9/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0120 - acc: 0.9959 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 10/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0118 - acc: 0.9961 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 11/20
60000/60000 [=====] - 9s 158us/step - loss: 0.0115 - acc: 0.9962 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 12/20
60000/60000 [=====] - 10s 159us/step - loss: 0.0117 - acc: 0.9961 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 13/20
60000/60000 [=====] - 9s 157us/step - loss: 0.0139 - acc: 0.9956 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 14/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0099 - acc: 0.9966 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 15/20

```

```

60000/60000 [=====] - 9s 155us/step - loss: 0.0064 - acc: 0.9979 - va
Epoch 16/20
60000/60000 [=====] - 9s 154us/step - loss: 0.0122 - acc: 0.9960 - va
Epoch 17/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0083 - acc: 0.9974 - va
Epoch 18/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0053 - acc: 0.9983 - va
Epoch 19/20
60000/60000 [=====] - 9s 158us/step - loss: 0.0088 - acc: 0.9974 - va
Epoch 20/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0086 - acc: 0.9974 - va

```

```

In [0]: score_relu2_adam_diff = model_relu.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu2_adam_diff[0])
        print('Test accuracy:', score_relu2_adam_diff[1])

        fig,ax = plt.subplots(1,1)
        ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

        # list of epoch numbers
        x = list(range(1,nb_epoch+1))

        # print(history.history.keys())
        # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
        # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

        # we will get val_loss and val_acc only when you pass the paramter validation_data
        # val_loss : validation loss
        # val_acc : validation accuracy

        # loss : training loss
        # acc : train accuracy
        # for each key in history.history we will have a list of length equal to number of ep


        vy = history.history['val_loss']
        ty = history.history['loss']
        plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.08052913746033337
Test accuracy: 0.9845

```

```
In [0]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)

```



3 MLP + ReLU + ADAM with 3 hidden layers

```

In [0]: model_relu = Sequential()

        model_relu.add(Dense(532, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
        model_relu.add(Dense(291, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
        model_relu.add(Dense(187, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
        model_relu.add(Dense(output_dim, activation='softmax'))

        print(model_relu.summary())

```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 532)	417620

```

-----
dense_9 (Dense)                (None, 291)                155103
-----
dense_10 (Dense)               (None, 187)                54604
-----
dense_11 (Dense)               (None, 10)                 1880
=====
Total params: 629,207
Trainable params: 629,207
Non-trainable params: 0
-----
None

```

```
In [0]: model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)
```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 9s 143us/step - loss: 0.2101 - acc: 0.9357 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 2/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0799 - acc: 0.9754 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 3/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0521 - acc: 0.9835 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 4/20
60000/60000 [=====] - 8s 138us/step - loss: 0.0385 - acc: 0.9880 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 5/20
60000/60000 [=====] - 8s 138us/step - loss: 0.0320 - acc: 0.9888 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 6/20
60000/60000 [=====] - 8s 140us/step - loss: 0.0241 - acc: 0.9918 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 7/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0238 - acc: 0.9922 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 8/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0209 - acc: 0.9935 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 9/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0184 - acc: 0.9941 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 10/20
60000/60000 [=====] - 8s 141us/step - loss: 0.0167 - acc: 0.9948 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 11/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0181 - acc: 0.9938 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 12/20
60000/60000 [=====] - 8s 138us/step - loss: 0.0145 - acc: 0.9955 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 13/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0141 - acc: 0.9956 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 14/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0156 - acc: 0.9951 - val_loss: 0.1875 - val_acc: 0.9450
Epoch 15/20

```

```

60000/60000 [=====] - 9s 147us/step - loss: 0.0103 - acc: 0.9967 - va
Epoch 16/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0111 - acc: 0.9963 - va
Epoch 17/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0178 - acc: 0.9944 - va
Epoch 18/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0082 - acc: 0.9974 - va
Epoch 19/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0089 - acc: 0.9971 - va
Epoch 20/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0106 - acc: 0.9969 - va

```

```

In [0]: score_relu3_adam = model_relu.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu3_adam[0])
        print('Test accuracy:', score_relu3_adam[1])

        fig,ax = plt.subplots(1,1)
        ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

        # list of epoch numbers
        x = list(range(1,nb_epoch+1))

        # print(history.history.keys())
        # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
        # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

        # we will get val_loss and val_acc only when you pass the paramter validation_data
        # val_loss : validation loss
        # val_acc : validation accuracy

        # loss : training loss
        # acc : train accuracy
        # for each key in history.history we will have a list of length equal to number of ep


        vy = history.history['val_loss']
        ty = history.history['loss']
        plt_dynamic(x, vy, ty, ax)

```

```

Test score: 0.10949685554472845
Test accuracy: 0.9773

```



```
In [0]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')
```

```
plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)
```



4 MLP + ReLU + ADAM with 5 hidden layers

```
In [0]: model_relu = Sequential()

model_relu.add(Dense(532, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0,
model_relu.add(Dense(443, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
model_relu.add(Dense(291, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
model_relu.add(Dense(167, activation='relu', kernel_initializer=RandomNormal(mean=0.0,
```

```

model_relu.add(Dense(125, activation='relu', kernel_initializer=RandomNormal(mean=0.0,

model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

```

```

-----
Layer (type)                 Output Shape              Param #
=====
dense_12 (Dense)             (None, 532)               417620
-----
dense_13 (Dense)             (None, 443)               236119
-----
dense_14 (Dense)             (None, 291)               129204
-----
dense_15 (Dense)             (None, 167)               48764
-----
dense_16 (Dense)             (None, 125)               21000
-----
dense_17 (Dense)             (None, 10)                1260
=====
Total params: 853,967
Trainable params: 853,967
Non-trainable params: 0
-----
None

```

```

In [0]: model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

```

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 12s 208us/step - loss: 0.2453 - acc: 0.9270 - val_loss: 0.0974 - val_acc: 0.9696

Epoch 2/20

60000/60000 [=====] - 12s 200us/step - loss: 0.0974 - acc: 0.9696 - val_loss: 0.0683 - val_acc: 0.9781

Epoch 3/20

60000/60000 [=====] - 12s 199us/step - loss: 0.0683 - acc: 0.9781 - val_loss: 0.0529 - val_acc: 0.9826

Epoch 4/20

60000/60000 [=====] - 12s 203us/step - loss: 0.0529 - acc: 0.9826 - val_loss: 0.0419 - val_acc: 0.9868

Epoch 5/20

60000/60000 [=====] - 12s 202us/step - loss: 0.0419 - acc: 0.9868 - val_loss: 0.0391 - val_acc: 0.9875

Epoch 6/20

60000/60000 [=====] - 12s 204us/step - loss: 0.0391 - acc: 0.9875 - val_loss: 0.0338 - val_acc: 0.9890

Epoch 7/20

60000/60000 [=====] - 12s 201us/step - loss: 0.0338 - acc: 0.9890 - val_loss: 0.0338 - val_acc: 0.9890

```

Epoch 8/20
60000/60000 [=====] - 12s 204us/step - loss: 0.0334 - acc: 0.9895 - va
Epoch 9/20
60000/60000 [=====] - 13s 210us/step - loss: 0.0263 - acc: 0.9920 - va
Epoch 10/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0273 - acc: 0.9918 - va
Epoch 11/20
60000/60000 [=====] - 12s 203us/step - loss: 0.0253 - acc: 0.9921 - va
Epoch 12/20
60000/60000 [=====] - 12s 203us/step - loss: 0.0219 - acc: 0.9928 - va
Epoch 13/20
60000/60000 [=====] - 12s 201us/step - loss: 0.0212 - acc: 0.9934 - va
Epoch 14/20
60000/60000 [=====] - 12s 202us/step - loss: 0.0172 - acc: 0.9946 - va
Epoch 15/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0195 - acc: 0.9938 - va
Epoch 16/20
60000/60000 [=====] - 12s 205us/step - loss: 0.0149 - acc: 0.9953 - va
Epoch 17/20
60000/60000 [=====] - 12s 202us/step - loss: 0.0140 - acc: 0.9957 - va
Epoch 18/20
60000/60000 [=====] - 12s 200us/step - loss: 0.0146 - acc: 0.9957 - va
Epoch 19/20
60000/60000 [=====] - 12s 197us/step - loss: 0.0141 - acc: 0.9960 - va
Epoch 20/20
60000/60000 [=====] - 12s 196us/step - loss: 0.0156 - acc: 0.9954 - va

```

```

In [0]: score_relu5_adam = model_relu.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu5_adam[0])
        print('Test accuracy:', score_relu5_adam[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy


```


for each key in history.history we will have a list of length equal to number of epochs

```
vy = history.history['val_loss']  
ty = history.history['loss']  
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1017726518217185

Test accuracy: 0.9787



```
In [0]: w_after = model_relu.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)  
h2_w = w_after[2].flatten().reshape(-1,1)  
h3_w = w_after[4].flatten().reshape(-1,1)  
h4_w = w_after[6].flatten().reshape(-1,1)  
h5_w = w_after[4].flatten().reshape(-1,1)  
out_w = w_after[10].flatten().reshape(-1,1)
```

```
fig = plt.figure()  
plt.title("Weight matrices after model trained")  
plt.subplot(1, 6, 1)  
plt.title("Trained model Weights")  
ax = sns.violinplot(y=h1_w,color='b')
```

```

plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='c')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='w')
plt.xlabel('Hidden Layer 5 ')

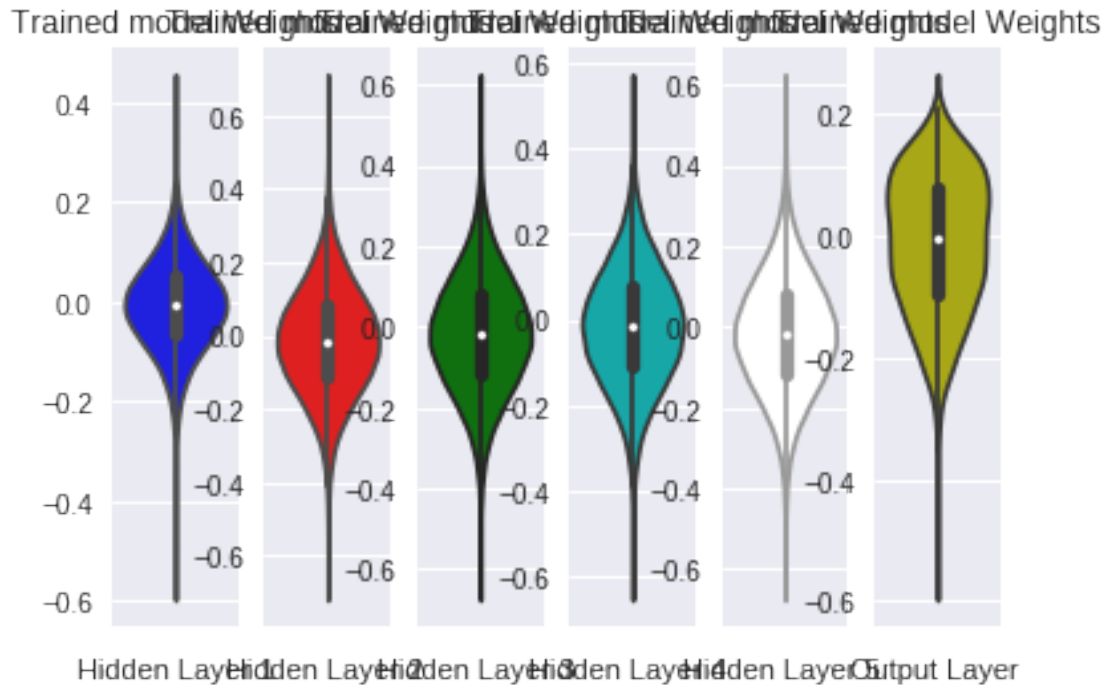
plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)

```



Observations 1. We used MLP with 2, 3 and 5 different layer and got accuracy ~98% but the difference between train loss and validation loss is slightly high and could so happen that we are overfitting. 2. Weight distributions is not too small and not too large and mean is at 0, which is a good sign of not getting into problem of vanishing or exploding gradient.

5 2. MLP + Batch-Norm + AdamOptimizer with 2 hidden layer

In [0]: *# Multilayer perceptron*

```
# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution N(0,) we satisfy this condition with
# h1 => =(2/(ni+ni+1) = 0.039 => N(0,) = N(0,0.039)
# h2 => =(2/(ni+ni+1) = 0.055 => N(0,) = N(0,0.055)
# h1 => =(2/(ni+ni+1) = 0.120 => N(0,) = N(0,0.120)
```

```
from keras.layers.normalization import BatchNormalization
```

```
model_batch = Sequential()
```

```
model_batch.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=2)))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.055, seed=2)))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(output_dim, activation='softmax'))
```

```
model_batch.summary()
```

```
-----
Layer (type)                 Output Shape          Param #
=====
dense_18 (Dense)             (None, 512)           401920
-----
batch_normalization_1 (Batch Normalization) (None, 512)           2048
-----
dense_19 (Dense)             (None, 128)           65664
-----
batch_normalization_2 (Batch Normalization) (None, 128)           512
-----
dense_20 (Dense)             (None, 10)            1290
=====
Total params: 471,434
Trainable params: 470,154
Non-trainable params: 1,280
-----
```

```
In [0]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, validation_data=(X_val, Y_val))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 138us/step - loss: 0.2004 - acc: 0.9403 - val_loss: 0.1000 - val_acc: 0.9700

Epoch 2/20

60000/60000 [=====] - 8s 130us/step - loss: 0.0724 - acc: 0.9788 - val_loss: 0.0319 - val_acc: 0.9903

Epoch 3/20

60000/60000 [=====] - 8s 127us/step - loss: 0.0470 - acc: 0.9858 - val_loss: 0.0206 - val_acc: 0.9934

Epoch 4/20

60000/60000 [=====] - 8s 126us/step - loss: 0.0319 - acc: 0.9903 - val_loss: 0.0252 - val_acc: 0.9921

Epoch 5/20

60000/60000 [=====] - 8s 128us/step - loss: 0.0252 - acc: 0.9921 - val_loss: 0.0161 - val_acc: 0.9952

Epoch 6/20

60000/60000 [=====] - 7s 125us/step - loss: 0.0206 - acc: 0.9934 - val_loss: 0.0160 - val_acc: 0.9951

Epoch 7/20

60000/60000 [=====] - 8s 127us/step - loss: 0.0161 - acc: 0.9952 - val_loss: 0.0124 - val_acc: 0.9962

Epoch 8/20

60000/60000 [=====] - 8s 127us/step - loss: 0.0160 - acc: 0.9951 - val_loss: 0.0124 - val_acc: 0.9962

Epoch 9/20

60000/60000 [=====] - 8s 126us/step - loss: 0.0124 - acc: 0.9962 - val_loss: 0.0124 - val_acc: 0.9962

Epoch 10/20

```

60000/60000 [=====] - 8s 125us/step - loss: 0.0113 - acc: 0.9966 - va
Epoch 11/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0115 - acc: 0.9965 - va
Epoch 12/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0081 - acc: 0.9978 - va
Epoch 13/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0099 - acc: 0.9968 - va
Epoch 14/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0067 - acc: 0.9978 - va
Epoch 15/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0134 - acc: 0.9953 - va
Epoch 16/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0094 - acc: 0.9968 - va
Epoch 17/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0067 - acc: 0.9976 - va
Epoch 18/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0055 - acc: 0.9983 - va
Epoch 19/20
60000/60000 [=====] - 8s 125us/step - loss: 0.0062 - acc: 0.9981 - va
Epoch 20/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0062 - acc: 0.9981 - va

```

```

In [0]: score_relu_bn2 = model_batch.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu_bn2[0])
        print('Test accuracy:', score_relu_bn2[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of e

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08353556353475215

Test accuracy: 0.9796



```
In [0]: w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
```

```
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)
```



6 MLP + Batch-Norm + AdamOptimizer with 2 hidden layer

```
In [0]: from keras.layers.normalization import BatchNormalization
```

```
model = Sequential()
```

```
model.add(Dense(526, activation = 'relu', input_shape=(input_dim,), kernel_initializer=
model_batch.add(BatchNormalization())
```

```
model.add(Dense(207, activation = 'relu', kernel_initializer = RandomNormal(mean = 0.0
model.add(BatchNormalization())
```

```
model.add(Dense(output_dim, activation = 'softmax'))
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
dense_21 (Dense)	(None, 526)	412910
dense_22 (Dense)	(None, 207)	109089
batch_normalization_4 (Batch Normalization)	(None, 207)	828
dense_23 (Dense)	(None, 10)	2080
Total params: 524,907		
Trainable params: 524,493		
Non-trainable params: 414		

```
In [0]: model_batch.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epochs)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 151us/step - loss: 6.3680 - acc: 0.9433 - val_loss: 4.5023 - val_acc: 0.9560

Epoch 2/20

60000/60000 [=====] - 8s 138us/step - loss: 1.6970 - acc: 0.9770 - val_loss: 4.5023 - val_acc: 0.9560

Epoch 3/20

60000/60000 [=====] - 8s 129us/step - loss: 4.0755 - acc: 0.9641 - val_loss: 4.5023 - val_acc: 0.9560

Epoch 4/20

60000/60000 [=====] - 8s 127us/step - loss: 2.9868 - acc: 0.9269 - val_loss: 4.5023 - val_acc: 0.9560

Epoch 5/20

60000/60000 [=====] - 8s 127us/step - loss: 3.3244 - acc: 0.9279 - val_loss: 4.5023 - val_acc: 0.9560

Epoch 6/20

60000/60000 [=====] - 8s 129us/step - loss: 3.1933 - acc: 0.9081 - val_loss: 4.5023 - val_acc: 0.9560

Epoch 7/20

60000/60000 [=====] - 8s 129us/step - loss: 3.7201 - acc: 0.9356 - val_loss: 4.5023 - val_acc: 0.9560

Epoch 8/20

60000/60000 [=====] - 8s 128us/step - loss: 4.5023 - acc: 0.9560 - val_loss: 4.5023 - val_acc: 0.9560

Epoch 9/20

60000/60000 [=====] - 8s 130us/step - loss: 4.5078 - acc: 0.9594 - val_loss: 4.5023 - val_acc: 0.9560

Epoch 10/20

60000/60000 [=====] - 8s 130us/step - loss: 4.6388 - acc: 0.9547 - val_loss: 4.5023 - val_acc: 0.9560

Epoch 11/20

60000/60000 [=====] - 8s 128us/step - loss: 2.7944 - acc: 0.9380 - val_loss: 4.5023 - val_acc: 0.9560

Epoch 12/20


```

60000/60000 [=====] - 8s 128us/step - loss: 4.2741 - acc: 0.9495 - va
Epoch 13/20
60000/60000 [=====] - 8s 127us/step - loss: 3.6275 - acc: 0.8936 - va
Epoch 14/20
60000/60000 [=====] - 8s 128us/step - loss: 4.8827 - acc: 0.9140 - va
Epoch 15/20
60000/60000 [=====] - 8s 127us/step - loss: 5.5093 - acc: 0.8756 - va
Epoch 16/20
60000/60000 [=====] - 8s 127us/step - loss: 6.0189 - acc: 0.8542 - va
Epoch 17/20
60000/60000 [=====] - 8s 127us/step - loss: 6.8726 - acc: 0.8445 - va
Epoch 18/20
60000/60000 [=====] - 8s 128us/step - loss: 7.4012 - acc: 0.8366 - va
Epoch 19/20
60000/60000 [=====] - 8s 128us/step - loss: 5.9204 - acc: 0.7755 - va
Epoch 20/20
60000/60000 [=====] - 8s 127us/step - loss: 6.1215 - acc: 0.7644 - va

```

```

In [0]: score_relu_bn2_diff = model_batch.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu_bn2_diff[0])
        print('Test accuracy:', score_relu_bn2_diff[1])

        fig,ax = plt.subplots(1,1)
        ax.set_xlabel('epoch')
        ax.set_ylabel('Categorical Crossentropy Loss')

        # list of epoch numbers
        x = list(range(1, nb_epoch+1))

        # print(history.history.keys())
        # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
        # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

        # we will get val_loss and val_acc only when you pass the paramter validation_data
        # val_loss : validation loss
        # val_acc : validation accuracy

        # loss : training loss
        # acc : train accuracy
        # for each key in history.history we will have a list of length equal to number of ep


        vy = history.history['val_loss']
        ty = history.history['loss']
        plt_dynamic(x, vy, ty, ax)

```

```

Test score: 6.459873273468018
Test accuracy: 0.8002

```



```
In [0]: w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)

```



7 MLP + Batch-Norm + AdamOptimizer with 3 hidden layer

```
In [0]: from keras.layers.normalization import BatchNormalization
```

```
model_batch = Sequential()
```

```
model_batch.add(Dense(345, activation = 'relu', input_shape=(input_dim,), kernel_initializer = RandomNormal(mean = 0, std = 0.01)))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(198, activation = 'relu', kernel_initializer = RandomNormal(mean = 0, std = 0.01)))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(57, activation = 'relu', kernel_initializer = RandomNormal(mean = 0, std = 0.01)))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(output_dim, activation = 'softmax'))
```

```
model_batch.summary()
```

```
-----
Layer (type)                 Output Shape              Param #
=====
dense_24 (Dense)             (None, 345)              270825
-----
batch_normalization_5 (Batch Normalization) (None, 345)              1380
-----
dense_25 (Dense)             (None, 198)              68508
-----
batch_normalization_6 (Batch Normalization) (None, 198)              792
-----
dense_26 (Dense)             (None, 57)               11343
-----
batch_normalization_7 (Batch Normalization) (None, 57)               228
-----
dense_27 (Dense)             (None, 10)               580
=====
Total params: 353,656
Trainable params: 352,456
Non-trainable params: 1,200
-----
```

```
In [0]: model_batch.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epochs)
```

```
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/20
60000/60000 [=====] - 8s 133us/step - loss: 0.2364 - acc: 0.9325 - val_loss: 0.0867 - val_acc: 0.9745
Epoch 2/20
60000/60000 [=====] - 7s 110us/step - loss: 0.0867 - acc: 0.9745 - val_loss: 0.0532 - val_acc: 0.9839
Epoch 3/20
60000/60000 [=====] - 7s 110us/step - loss: 0.0532 - acc: 0.9839 - val_loss: 0.0395 - val_acc: 0.9875
Epoch 4/20
60000/60000 [=====] - 7s 111us/step - loss: 0.0395 - acc: 0.9875 - val_loss: 0.0286 - val_acc: 0.9913
Epoch 5/20
60000/60000 [=====] - 7s 110us/step - loss: 0.0286 - acc: 0.9913 - val_loss: 0.0237 - val_acc: 0.9926
Epoch 6/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0237 - acc: 0.9926 - val_loss: 0.0179 - val_acc: 0.9944
Epoch 7/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0179 - acc: 0.9944 - val_loss: 0.0184 - val_acc: 0.9940
Epoch 8/20
60000/60000 [=====] - 7s 111us/step - loss: 0.0184 - acc: 0.9940 - val_loss: 0.0184 - val_acc: 0.9940
Epoch 9/20
```

```

60000/60000 [=====] - 7s 111us/step - loss: 0.0157 - acc: 0.9949 - va
Epoch 10/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0133 - acc: 0.9957 - va
Epoch 11/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0151 - acc: 0.9951 - va
Epoch 12/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0121 - acc: 0.9960 - va
Epoch 13/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0129 - acc: 0.9958 - va
Epoch 14/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0100 - acc: 0.9969 - va
Epoch 15/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0091 - acc: 0.9972 - va
Epoch 16/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0096 - acc: 0.9969 - va
Epoch 17/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0084 - acc: 0.9970 - va
Epoch 18/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0078 - acc: 0.9974 - va
Epoch 19/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0074 - acc: 0.9977 - va
Epoch 20/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0057 - acc: 0.9982 - va

```

```

In [0]: score_relu_bn3 = model_batch.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu_bn3[0])
        print('Test accuracy:', score_relu_bn3[1])

        fig,ax = plt.subplots(1,1)
        ax.set_xlabel('epoch')
        ax.set_ylabel('Categorical Crossentropy Loss')

        # list of epoch numbers
        x = list(range(1, nb_epoch+1))

        # print(history.history.keys())
        # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
        # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

        # we will get val_loss and val_acc only when you pass the paramter validation_data
        # val_loss : validation loss
        # val_acc : validation accuracy

        # loss : training loss
        # acc : train accuracy
        # for each key in history.history we will have a list of length equal to number of ep

```

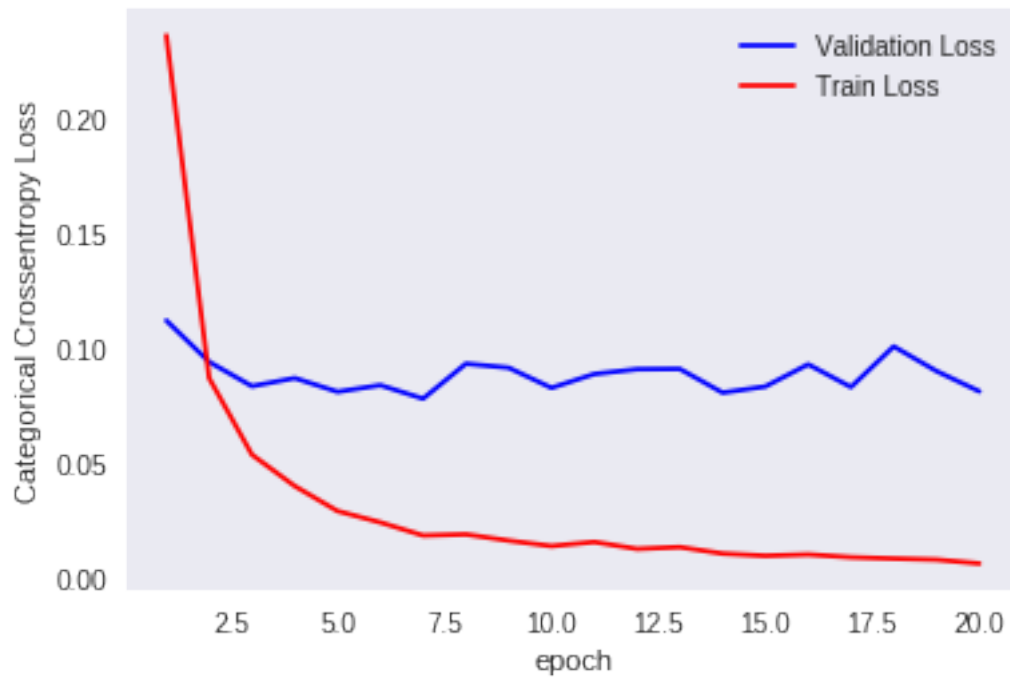
```

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08087161136912183

Test accuracy: 0.9798



```

In [0]: w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1, 1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")

```

```

ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')


plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)

```



8 MLP + Batch-Norm + AdamOptimizer with 5 hidden layer

```
In [0]: from keras.layers.normalization import BatchNormalization
```

```

model_batch = Sequential()

model_batch.add(Dense(451, activation = 'relu', input_shape=(input_dim,), kernel_initializer = RandomNormal(mean=0.0, stddev=0.01)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(272, activation = 'relu', kernel_initializer = RandomNormal(mean=0.0, stddev=0.01)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(180, activation = 'relu', kernel_initializer = RandomNormal(mean=0.0, stddev=0.01)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(97, activation = 'relu', kernel_initializer = RandomNormal(mean=0.0, stddev=0.01)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(56, activation = 'relu', kernel_initializer = RandomNormal(mean=0.0, stddev=0.01)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation = 'softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
dense_28 (Dense)	(None, 451)	354035
batch_normalization_8 (Batch Normalization)	(None, 451)	1804
dense_29 (Dense)	(None, 272)	122944
batch_normalization_9 (Batch Normalization)	(None, 272)	1088
dense_30 (Dense)	(None, 180)	49140
batch_normalization_10 (Batch Normalization)	(None, 180)	720
dense_31 (Dense)	(None, 97)	17557
batch_normalization_11 (Batch Normalization)	(None, 97)	388
dense_32 (Dense)	(None, 56)	5488
batch_normalization_12 (Batch Normalization)	(None, 56)	224
dense_33 (Dense)	(None, 10)	570

Total params: 553,958
Trainable params: 551,846
Non-trainable params: 2,112

```
-----  
In [0]: model_batch.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = [  
        history = model_batch.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20  
60000/60000 [=====] - 13s 209us/step - loss: 0.2851 - acc: 0.9181 - va  
Epoch 2/20  
60000/60000 [=====] - 11s 182us/step - loss: 0.0984 - acc: 0.9705 - va  
Epoch 3/20  
60000/60000 [=====] - 11s 179us/step - loss: 0.0643 - acc: 0.9797 - va  
Epoch 4/20  
60000/60000 [=====] - 11s 178us/step - loss: 0.0470 - acc: 0.9852 - va  
Epoch 5/20  
60000/60000 [=====] - 11s 178us/step - loss: 0.0364 - acc: 0.9880 - va  
Epoch 6/20  
60000/60000 [=====] - 11s 178us/step - loss: 0.0343 - acc: 0.9886 - va  
Epoch 7/20  
60000/60000 [=====] - 10s 172us/step - loss: 0.0266 - acc: 0.9913 - va  
Epoch 8/20  
60000/60000 [=====] - 10s 173us/step - loss: 0.0242 - acc: 0.9919 - va  
Epoch 9/20  
60000/60000 [=====] - 10s 174us/step - loss: 0.0226 - acc: 0.9927 - va  
Epoch 10/20  
60000/60000 [=====] - 10s 173us/step - loss: 0.0201 - acc: 0.9934 - va  
Epoch 11/20  
60000/60000 [=====] - 10s 172us/step - loss: 0.0183 - acc: 0.9940 - va  
Epoch 12/20  
60000/60000 [=====] - 10s 174us/step - loss: 0.0180 - acc: 0.9940 - va  
Epoch 13/20  
60000/60000 [=====] - 10s 175us/step - loss: 0.0161 - acc: 0.9947 - va  
Epoch 14/20  
60000/60000 [=====] - 10s 172us/step - loss: 0.0135 - acc: 0.9956 - va  
Epoch 15/20  
60000/60000 [=====] - 10s 173us/step - loss: 0.0137 - acc: 0.9956 - va  
Epoch 16/20  
60000/60000 [=====] - 11s 180us/step - loss: 0.0110 - acc: 0.9965 - va  
Epoch 17/20  
60000/60000 [=====] - 11s 180us/step - loss: 0.0113 - acc: 0.9964 - va  
Epoch 18/20  
60000/60000 [=====] - 11s 180us/step - loss: 0.0140 - acc: 0.9953 - va  
Epoch 19/20
```

```
60000/60000 [=====] - 11s 181us/step - loss: 0.0108 - acc: 0.9965 - v
Epoch 20/20
60000/60000 [=====] - 11s 181us/step - loss: 0.0111 - acc: 0.9964 - v
```

```
In [0]: score_relu_bn5 = model_batch.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu_bn5[0])
        print('Test accuracy:', score_relu_bn5[1])

        fig,ax = plt.subplots(1,1)
        ax.set_xlabel('epoch')
        ax.set_ylabel('Categorical Crossentropy Loss')

        # list of epoch numbers
        x = list(range(1, nb_epoch+1))


        # print(history.history.keys())
        # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
        # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

        # we will get val_loss and val_acc only when you pass the paramter validation_data
        # val_loss : validation loss
        # val_acc : validation accuracy

        # loss : training loss
        # acc : train accuracy
        # for each key in history.history we will have a list of length equal to number of e

        vy = history.history['val_loss']
        ty = history.history['loss']
        plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.0770734781166655
Test accuracy: 0.9815
```



```
In [0]: w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
```

```

plt.xlabel('Hidden Layer 3')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='c')
plt.xlabel('Hidden Layer 4')


plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='w')
plt.xlabel('Hidden Layer 5')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer')
plt.show()

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is deprecated. Use kde_data = remove_na(group_data)

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is deprecated. Use violin_data = remove_na(group_data)



Observations 1. We did not get good accuracy with 2 and 3 hidden layer and also mean of weights deviated from 0 to 1. 2. With 5 hidden layer we got slightly good accuracy than 2 and 3 hidden layer.

9 3. MLP + Dropout + AdamOptimizer with 2 hidden layer

```
In [0]: from keras.layers import Dropout
```

```
model_drop = Sequential()
```

```
model_drop.add(Dense(370, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(112, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

```
-----
Layer (type)                 Output Shape          Param #
=====
dense_34 (Dense)              (None, 370)           290450
-----
dropout_1 (Dropout)           (None, 370)           0
-----
dense_35 (Dense)              (None, 112)           41552
-----
dropout_2 (Dropout)           (None, 112)           0
-----
dense_36 (Dense)              (None, 10)            1130
=====
Total params: 333,132
Trainable params: 333,132
Non-trainable params: 0
-----
```

```
In [0]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 116us/step - loss: 0.7979 - acc: 0.7623 - val_loss: 0.7979 - val_acc: 0.7623

Epoch 2/20

60000/60000 [=====] - 6s 94us/step - loss: 0.3699 - acc: 0.8913 - val_loss: 0.3699 - val_acc: 0.8913

Epoch 3/20

60000/60000 [=====] - 6s 95us/step - loss: 0.2955 - acc: 0.9159 - val_loss: 0.2955 - val_acc: 0.9159

Epoch 4/20

60000/60000 [=====] - 6s 96us/step - loss: 0.2479 - acc: 0.9279 - val_loss: 0.2479 - val_acc: 0.9279

```

Epoch 5/20
60000/60000 [=====] - 6s 100us/step - loss: 0.2196 - acc: 0.9374 - val
Epoch 6/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1952 - acc: 0.9437 - val
Epoch 7/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1912 - acc: 0.9458 - val
Epoch 8/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1760 - acc: 0.9496 - val
Epoch 9/20
60000/60000 [=====] - 6s 96us/step - loss: 0.1610 - acc: 0.9533 - val
Epoch 10/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1514 - acc: 0.9563 - val
Epoch 11/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1486 - acc: 0.9580 - val
Epoch 12/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1395 - acc: 0.9600 - val
Epoch 13/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1316 - acc: 0.9619 - val
Epoch 14/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1259 - acc: 0.9637 - val
Epoch 15/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1232 - acc: 0.9640 - val
Epoch 16/20
60000/60000 [=====] - 6s 96us/step - loss: 0.1162 - acc: 0.9666 - val
Epoch 17/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1171 - acc: 0.9664 - val
Epoch 18/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1115 - acc: 0.9670 - val
Epoch 19/20
60000/60000 [=====] - 6s 99us/step - loss: 0.1006 - acc: 0.9709 - val
Epoch 20/20
60000/60000 [=====] - 6s 102us/step - loss: 0.1003 - acc: 0.9707 - val

```

```

In [0]: score_relu_drop2 = model_drop.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu_drop2[0])
        print('Test accuracy:', score_relu_drop2[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

```

```

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0825668051331857

Test accuracy: 0.9783



```

In [0]: w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")

```

```

plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)

```



10 MLP + Dropout + AdamOptimizer with 3 hidden layer

```
In [0]: from keras.layers import Dropout
```

```
model_drop = Sequential()
```

```
model_drop.add(Dense(531, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(375, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(130, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_37 (Dense)	(None, 531)	416835
dropout_3 (Dropout)	(None, 531)	0
dense_38 (Dense)	(None, 375)	199500
dropout_4 (Dropout)	(None, 375)	0
dense_39 (Dense)	(None, 130)	48880
dropout_5 (Dropout)	(None, 130)	0
dense_40 (Dense)	(None, 10)	1310
Total params: 666,525		
Trainable params: 666,525		
Non-trainable params: 0		

```
In [0]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20

```

60000/60000 [=====] - 12s 192us/step - loss: 10.2589 - acc: 0.3486 - v
Epoch 2/20
60000/60000 [=====] - 10s 169us/step - loss: 7.3090 - acc: 0.5348 - va
Epoch 3/20
60000/60000 [=====] - 10s 170us/step - loss: 5.3924 - acc: 0.6534 - va
Epoch 4/20
60000/60000 [=====] - 10s 170us/step - loss: 4.1462 - acc: 0.7304 - va
Epoch 5/20
60000/60000 [=====] - 10s 170us/step - loss: 3.3093 - acc: 0.7818 - va
Epoch 6/20
60000/60000 [=====] - 10s 170us/step - loss: 2.6803 - acc: 0.8215 - va
Epoch 7/20
60000/60000 [=====] - 10s 171us/step - loss: 2.3632 - acc: 0.8423 - va
Epoch 8/20
60000/60000 [=====] - 10s 170us/step - loss: 2.1355 - acc: 0.8575 - va
Epoch 9/20
60000/60000 [=====] - 10s 173us/step - loss: 1.8977 - acc: 0.8719 - va
Epoch 10/20
60000/60000 [=====] - 10s 172us/step - loss: 1.7798 - acc: 0.8790 - va
Epoch 11/20
60000/60000 [=====] - 10s 171us/step - loss: 1.6786 - acc: 0.8857 - va
Epoch 12/20
60000/60000 [=====] - 10s 171us/step - loss: 1.5114 - acc: 0.8957 - va
Epoch 13/20
60000/60000 [=====] - 10s 172us/step - loss: 1.4467 - acc: 0.8977 - va
Epoch 14/20
60000/60000 [=====] - 10s 170us/step - loss: 1.2593 - acc: 0.9055 - va
Epoch 15/20
60000/60000 [=====] - 10s 170us/step - loss: 1.1119 - acc: 0.9063 - va
Epoch 16/20
60000/60000 [=====] - 10s 169us/step - loss: 0.9847 - acc: 0.9029 - va
Epoch 17/20
60000/60000 [=====] - 10s 168us/step - loss: 0.8042 - acc: 0.8968 - va
Epoch 18/20
60000/60000 [=====] - 10s 171us/step - loss: 0.6922 - acc: 0.8995 - va
Epoch 19/20
60000/60000 [=====] - 10s 168us/step - loss: 0.6329 - acc: 0.9009 - va
Epoch 20/20
60000/60000 [=====] - 10s 170us/step - loss: 0.5741 - acc: 0.9026 - va

```

```

In [0]: score_relu_drop3 = model_drop.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu_drop3[0])
        print('Test accuracy:', score_relu_drop3[1])

        fig,ax = plt.subplots(1,1)
        ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

```

```

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.37853595131413315

Test accuracy: 0.9328



```
In [0]: w_after = model_drop.get_weights()
```

```

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')


plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer')
plt.show()

```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)

```



11 MLP + Dropout + AdamOptimizer with 5 hidden layer

```
In [0]: from keras.layers import Dropout
```

```
model_drop = Sequential()
```

```
model_drop.add(Dense(681, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(475, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(230, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(102, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(28, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(Dropout(0.5))
```

```
model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_41 (Dense)	(None, 681)	534585
dropout_6 (Dropout)	(None, 681)	0
dense_42 (Dense)	(None, 475)	323950
dropout_7 (Dropout)	(None, 475)	0
dense_43 (Dense)	(None, 230)	109480
dropout_8 (Dropout)	(None, 230)	0
dense_44 (Dense)	(None, 102)	23562
dropout_9 (Dropout)	(None, 102)	0
dense_45 (Dense)	(None, 28)	2884
dropout_10 (Dropout)	(None, 28)	0
dense_46 (Dense)	(None, 10)	290
Total params: 994,751		
Trainable params: 994,751		
Non-trainable params: 0		

```
In [0]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 16s 274us/step - loss: 14.4528 - acc: 0.1023 - val_loss: 14.4528 - val_acc: 0.1023

Epoch 2/20

60000/60000 [=====] - 15s 249us/step - loss: 14.4791 - acc: 0.1008 - val_loss: 14.4791 - val_acc: 0.1008

Epoch 3/20

60000/60000 [=====] - 15s 250us/step - loss: 14.4756 - acc: 0.1010 - val_loss: 14.4756 - val_acc: 0.1010

Epoch 4/20

60000/60000 [=====] - 15s 250us/step - loss: 14.5060 - acc: 0.0991 - val_loss: 14.5060 - val_acc: 0.0991

Epoch 5/20

60000/60000 [=====] - 15s 249us/step - loss: 14.4631 - acc: 0.1014 - val_loss: 14.4631 - val_acc: 0.1014

Epoch 6/20

60000/60000 [=====] - 15s 253us/step - loss: 14.4937 - acc: 0.0995 - val_loss: 14.4937 - val_acc: 0.0995

```

Epoch 7/20
60000/60000 [=====] - 15s 252us/step - loss: 14.4749 - acc: 0.1003 - v
Epoch 8/20
60000/60000 [=====] - 15s 248us/step - loss: 14.5048 - acc: 0.0985 - v
Epoch 9/20
60000/60000 [=====] - 15s 246us/step - loss: 14.4717 - acc: 0.1007 - v
Epoch 10/20
60000/60000 [=====] - 15s 248us/step - loss: 14.4314 - acc: 0.1035 - v
Epoch 11/20
60000/60000 [=====] - 15s 246us/step - loss: 14.4519 - acc: 0.1023 - v
Epoch 12/20
60000/60000 [=====] - 15s 245us/step - loss: 14.4497 - acc: 0.1024 - v
Epoch 13/20
60000/60000 [=====] - 15s 244us/step - loss: 14.4248 - acc: 0.1043 - v
Epoch 14/20
60000/60000 [=====] - 15s 247us/step - loss: 14.4156 - acc: 0.1049 - v
Epoch 15/20
60000/60000 [=====] - 15s 245us/step - loss: 14.4187 - acc: 0.1047 - v
Epoch 16/20
60000/60000 [=====] - 15s 247us/step - loss: 14.4085 - acc: 0.1054 - v
Epoch 17/20
60000/60000 [=====] - 15s 249us/step - loss: 14.3898 - acc: 0.1066 - v
Epoch 18/20
60000/60000 [=====] - 15s 247us/step - loss: 14.3788 - acc: 0.1071 - v
Epoch 19/20
60000/60000 [=====] - 15s 249us/step - loss: 14.4300 - acc: 0.1036 - v
Epoch 20/20
60000/60000 [=====] - 15s 247us/step - loss: 14.4250 - acc: 0.1037 - v

```

```

In [0]: score_relu_drop5 = model_drop.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu_drop5[0])
        print('Test accuracy:', score_relu_drop5[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

```

```

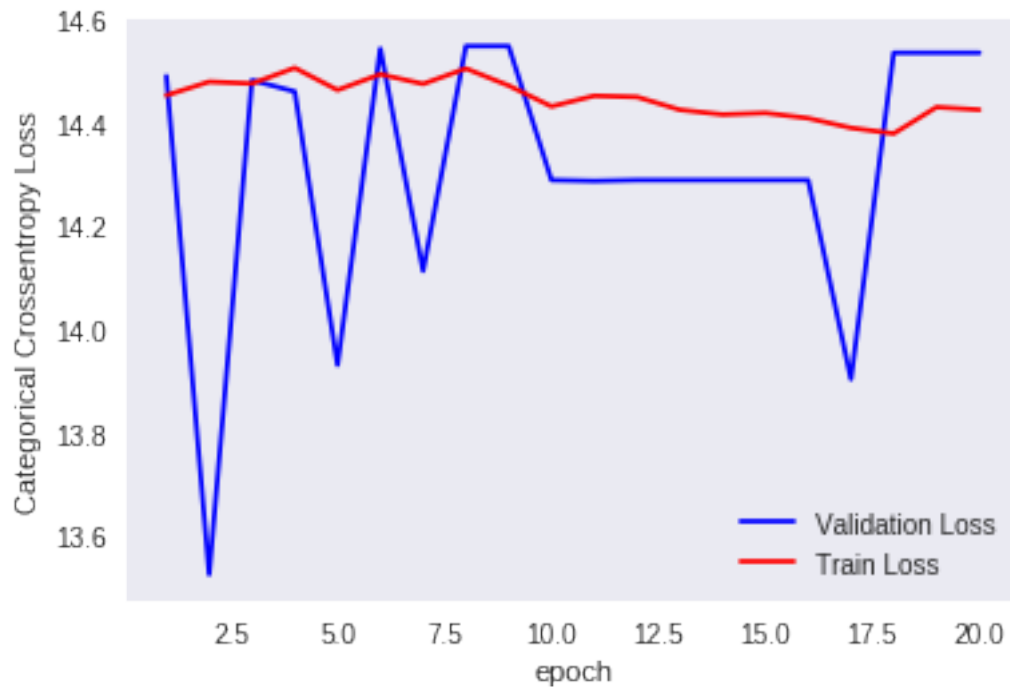
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 14.535298265075683

Test accuracy: 0.0982



```
In [0]: w_after = model_drop.get_weights()
```

```

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

```

```

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)

```



```

plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='c')
plt.xlabel('Hidden Layer 4')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='w')
plt.xlabel('Hidden Layer 5')

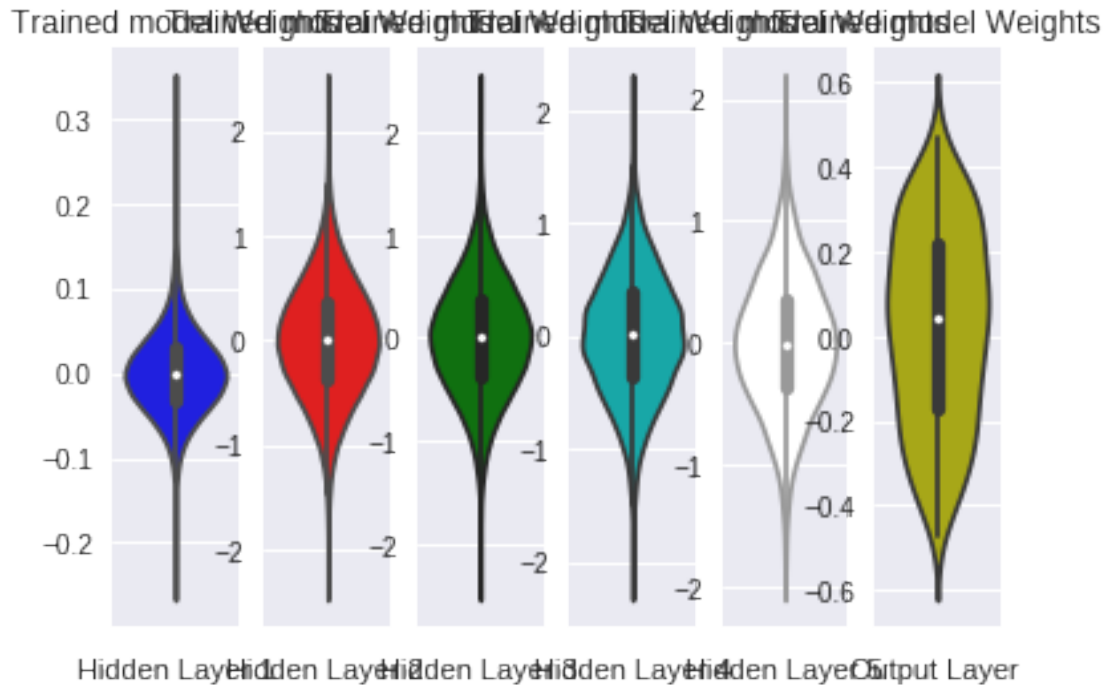
plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)

```



Observations 1. We use 2, 3 and 5 different layer architecture with dropout and did not get good accuracy in 5 layer architecture whereas first 2 model architecture works very well.

12 4. MLP + RELU + Dropout + BatchNormalization + AdamOptimizer with 2 hidden layers

In [0]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization->

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_47 (Dense)	(None, 512)	401920
batch_normalization_13 (Batch Normalization)	(None, 512)	2048
dropout_11 (Dropout)	(None, 512)	0
dense_48 (Dense)	(None, 128)	65664
batch_normalization_14 (Batch Normalization)	(None, 128)	512
dropout_12 (Dropout)	(None, 128)	0
dense_49 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

```
In [0]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 10s 169us/step - loss: 0.4788 - acc: 0.8538 - val_loss: 0.2481 - val_acc: 0.9256
Epoch 2/20
60000/60000 [=====] - 8s 139us/step - loss: 0.2481 - acc: 0.9256 - val_loss: 0.2007 - val_acc: 0.9398
Epoch 3/20
60000/60000 [=====] - 8s 139us/step - loss: 0.2007 - acc: 0.9398 - val_loss: 0.1705 - val_acc: 0.9483
Epoch 4/20
60000/60000 [=====] - 8s 141us/step - loss: 0.1705 - acc: 0.9483 - val_loss: 0.1532 - val_acc: 0.9536
Epoch 5/20
60000/60000 [=====] - 8s 141us/step - loss: 0.1532 - acc: 0.9536 - val_loss: 0.1423 - val_acc: 0.9571
Epoch 6/20
60000/60000 [=====] - 8s 139us/step - loss: 0.1423 - acc: 0.9571 - val_loss: 0.1290 - val_acc: 0.9613
Epoch 7/20
60000/60000 [=====] - 8s 139us/step - loss: 0.1290 - acc: 0.9613 - val_loss: 0.1207 - val_acc: 0.9627
Epoch 8/20
60000/60000 [=====] - 8s 138us/step - loss: 0.1207 - acc: 0.9627 - val_loss: 0.1142 - val_acc: 0.9656
Epoch 9/20
60000/60000 [=====] - 8s 138us/step - loss: 0.1142 - acc: 0.9656 - val_loss: 0.1076 - val_acc: 0.9674
Epoch 10/20
60000/60000 [=====] - 8s 139us/step - loss: 0.1076 - acc: 0.9674 - val_loss: 0.1076 - val_acc: 0.9674
```

```

Epoch 11/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0999 - acc: 0.9692 - va
Epoch 12/20
60000/60000 [=====] - 9s 145us/step - loss: 0.0941 - acc: 0.9710 - va
Epoch 13/20
60000/60000 [=====] - 8s 140us/step - loss: 0.0900 - acc: 0.9712 - va
Epoch 14/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0884 - acc: 0.9721 - va
Epoch 15/20
60000/60000 [=====] - 8s 140us/step - loss: 0.0829 - acc: 0.9741 - va
Epoch 16/20
60000/60000 [=====] - 8s 140us/step - loss: 0.0826 - acc: 0.9745 - va
Epoch 17/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0762 - acc: 0.9761 - va
Epoch 18/20
60000/60000 [=====] - 8s 140us/step - loss: 0.0724 - acc: 0.9775 - va
Epoch 19/20
60000/60000 [=====] - 8s 140us/step - loss: 0.0705 - acc: 0.9781 - va
Epoch 20/20
60000/60000 [=====] - 8s 137us/step - loss: 0.0669 - acc: 0.9792 - va

```

```

In [0]: score_relu_bn_drop2 = model_drop.evaluate(X_test, Y_test, verbose=0)
        print('Test score:', score_relu_bn_drop2[0])
        print('Test accuracy:', score_relu_bn_drop2[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of ep

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0575819099090164

Test accuracy: 0.9831



```
In [0]: w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Output Layer ')
plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)
```



13 MLP + RELU + Dropout + BatchNormalization with 2 hidden layers

```
In [0]: from keras.layers.normalization import BatchNormalization
        from keras.layers import Dropout
        from keras.models import Sequential
        from keras.layers import Dense, Activation
```

```
model = Sequential()
```

```
model.add(Dense(389, activation = "relu", input_shape = (input_dim,), kernel_initializer=
model.add(BatchNormalization())
model.add(Dropout(0.5))
```

```
model.add(Dense(258, activation = "relu", kernel_initializer = RandomNormal(mean=0.0, s
model.add(BatchNormalization())
```

```

model.add(Dropout(0.5))

model.add(Dense(output_dim, activation = "softmax"))

model.summary()

```

Layer (type)	Output Shape	Param #
dense_50 (Dense)	(None, 389)	305365
batch_normalization_15 (Batch Normalization)	(None, 389)	1556
dropout_13 (Dropout)	(None, 389)	0
dense_51 (Dense)	(None, 258)	100620
batch_normalization_16 (Batch Normalization)	(None, 258)	1032
dropout_14 (Dropout)	(None, 258)	0
dense_52 (Dense)	(None, 10)	2590
Total params: 411,163		
Trainable params: 409,869		
Non-trainable params: 1,294		

```
In [0]: model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])
```

```
history = model.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=0)
```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 10s 166us/step - loss: 0.4825 - acc: 0.8528 - val_loss: 0.4825 - val_acc: 0.8528
Epoch 2/20
60000/60000 [=====] - 8s 135us/step - loss: 0.2447 - acc: 0.9261 - val_loss: 0.2447 - val_acc: 0.9261
Epoch 3/20
60000/60000 [=====] - 8s 136us/step - loss: 0.1989 - acc: 0.9403 - val_loss: 0.1989 - val_acc: 0.9403
Epoch 4/20
60000/60000 [=====] - 8s 136us/step - loss: 0.1704 - acc: 0.9486 - val_loss: 0.1704 - val_acc: 0.9486
Epoch 5/20
60000/60000 [=====] - 8s 135us/step - loss: 0.1525 - acc: 0.9535 - val_loss: 0.1525 - val_acc: 0.9535
Epoch 6/20
60000/60000 [=====] - 8s 135us/step - loss: 0.1348 - acc: 0.9576 - val_loss: 0.1348 - val_acc: 0.9576
Epoch 7/20
60000/60000 [=====] - 8s 135us/step - loss: 0.1287 - acc: 0.9605 - val_loss: 0.1287 - val_acc: 0.9605

```

```

Epoch 8/20
60000/60000 [=====] - 8s 137us/step - loss: 0.1191 - acc: 0.9629 - va
Epoch 9/20
60000/60000 [=====] - 8s 137us/step - loss: 0.1101 - acc: 0.9658 - va
Epoch 10/20
60000/60000 [=====] - 8s 134us/step - loss: 0.1092 - acc: 0.9666 - va
Epoch 11/20
60000/60000 [=====] - 8s 133us/step - loss: 0.1044 - acc: 0.9681 - va
Epoch 12/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0975 - acc: 0.9693 - va
Epoch 13/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0937 - acc: 0.9706 - va
Epoch 14/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0924 - acc: 0.9707 - va
Epoch 15/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0855 - acc: 0.9734 - va
Epoch 16/20
60000/60000 [=====] - 8s 136us/step - loss: 0.0830 - acc: 0.9734 - va
Epoch 17/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0784 - acc: 0.9753 - va
Epoch 18/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0798 - acc: 0.9753 - va
Epoch 19/20
60000/60000 [=====] - 8s 136us/step - loss: 0.0733 - acc: 0.9763 - va
Epoch 20/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0710 - acc: 0.9775 - va

```

```

In [0]: score_relu_bn_drop2_diff = model.evaluate(X_test, Y_test, verbose = 0)
        print('Test score:', score_relu_bn_drop2_diff[0])
        print('Test accuracy:', score_relu_bn_drop2_diff[1])

        fig, ax = plt.subplots(1, 1)
        ax.set_xlabel('epoch')
        ax.set_ylabel('Categorical Crossentropy Loss')

        # list of epoch numbers
        x = list(range(1,nb_epoch+1))

        # print(history.history.keys())
        # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
        # history = model_drop2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

        # we will get val_loss and val_acc only when you pass the paramter validation_data
        # val_loss : validation loss
        # val_acc : validation accuracy

        # loss : training loss

```



```


# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06233537118713721

Test accuracy: 0.982



```
In [0]: w_after = model.get_weights()
```

```

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

```

```


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
violin_data = remove_na(group_data)
```



14 MLP + RELU + Dropout + BatchNormalization with 3 hidden layers

```
In [0]: from keras.layers.normalization import BatchNormalization
        from keras.layers import Dropout
        from keras.models import Sequential
        from keras.layers import Dense, Activation
```

```

model = Sequential()

model.add(Dense(434, activation = "relu", input_shape = (input_dim,), kernel_initializer=
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(391, activation = "relu", kernel_initializer = RandomNormal(mean=0.0, s
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(141, activation = "relu", kernel_initializer = RandomNormal(mean=0.0, s
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(output_dim, activation = "softmax"))

model.summary()

```

```

-----
Layer (type)                 Output Shape              Param #
=====
dense_53 (Dense)             (None, 434)              340690
-----
batch_normalization_17 (Batc (None, 434)              1736
-----
dropout_15 (Dropout)         (None, 434)              0
-----
dense_54 (Dense)             (None, 391)              170085
-----
batch_normalization_18 (Batc (None, 391)              1564
-----
dropout_16 (Dropout)         (None, 391)              0
-----
dense_55 (Dense)             (None, 141)              55272
-----
batch_normalization_19 (Batc (None, 141)              564
-----
dropout_17 (Dropout)         (None, 141)              0
-----
dense_56 (Dense)             (None, 10)               1420
=====
Total params: 571,331
Trainable params: 569,399
Non-trainable params: 1,932
-----

```

```

In [0]: model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])

```

```

        history = model.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=0)

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 13s 224us/step - loss: 0.7627 - acc: 0.7634 - va
Epoch 2/20
60000/60000 [=====] - 11s 185us/step - loss: 0.3444 - acc: 0.8963 - va
Epoch 3/20
60000/60000 [=====] - 11s 185us/step - loss: 0.2640 - acc: 0.9214 - va
Epoch 4/20
60000/60000 [=====] - 11s 185us/step - loss: 0.2262 - acc: 0.9333 - va
Epoch 5/20
60000/60000 [=====] - 11s 186us/step - loss: 0.1982 - acc: 0.9404 - va
Epoch 6/20
60000/60000 [=====] - 12s 194us/step - loss: 0.1781 - acc: 0.9472 - va
Epoch 7/20
60000/60000 [=====] - 11s 185us/step - loss: 0.1633 - acc: 0.9515 - va
Epoch 8/20
60000/60000 [=====] - 11s 186us/step - loss: 0.1464 - acc: 0.9557 - va
Epoch 9/20
60000/60000 [=====] - 11s 187us/step - loss: 0.1406 - acc: 0.9583 - va
Epoch 10/20
60000/60000 [=====] - 11s 187us/step - loss: 0.1367 - acc: 0.9592 - va
Epoch 11/20
60000/60000 [=====] - 11s 186us/step - loss: 0.1256 - acc: 0.9618 - va
Epoch 12/20
60000/60000 [=====] - 11s 186us/step - loss: 0.1205 - acc: 0.9647 - va
Epoch 13/20
60000/60000 [=====] - 11s 186us/step - loss: 0.1155 - acc: 0.9656 - va
Epoch 14/20
60000/60000 [=====] - 11s 186us/step - loss: 0.1110 - acc: 0.9672 - va
Epoch 15/20
60000/60000 [=====] - 11s 186us/step - loss: 0.1054 - acc: 0.9677 - va
Epoch 16/20
60000/60000 [=====] - 11s 189us/step - loss: 0.1017 - acc: 0.9691 - va
Epoch 17/20
60000/60000 [=====] - 11s 187us/step - loss: 0.0962 - acc: 0.9712 - va
Epoch 18/20
60000/60000 [=====] - 11s 187us/step - loss: 0.0937 - acc: 0.9725 - va
Epoch 19/20
60000/60000 [=====] - 11s 187us/step - loss: 0.0873 - acc: 0.9729 - va
Epoch 20/20
60000/60000 [=====] - 11s 187us/step - loss: 0.0884 - acc: 0.9729 - va

In [0]: score_relu_bn_drop3 = model.evaluate(X_test, Y_test, verbose = 0)
        print('Test score:', score_relu_bn_drop3[0])

```

```

print('Test accuracy:', score_relu_bn_drop3[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of e

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0605601929532073

Test accuracy: 0.9824



```

In [0]: w_after = model.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer')
plt.show()

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)

```



15 MLP + RELU + Dropout + BatchNormalization with 5 hidden layers

```
In [0]: model = Sequential()
```

```
model.add(Dense(697, activation = "relu", input_shape = (input_dim,), kernel_initializer=RandomNormal(mean=0.0, std=0.01))
model.add(BatchNormalization())
model.add(Dropout(0.5))
```

```
model.add(Dense(458, activation = "relu", kernel_initializer = RandomNormal(mean=0.0, std=0.01))
model.add(BatchNormalization())
model.add(Dropout(0.5))
```

```
model.add(Dense(246, activation = "relu", kernel_initializer = RandomNormal(mean=0.0, std=0.01))
model.add(BatchNormalization())
model.add(Dropout(0.5))
```

```
model.add(Dense(111, activation = "relu", kernel_initializer = RandomNormal(mean=0.0, std=0.01))
model.add(BatchNormalization())
model.add(Dropout(0.5))
```

```
model.add(Dense(58, activation = "relu", kernel_initializer = RandomNormal(mean=0.0, std=0.01))
model.add(BatchNormalization())
model.add(Dropout(0.5))
```

```
model.add(Dense(output_dim, activation = "softmax"))
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
dense_57 (Dense)	(None, 697)	547145
batch_normalization_20 (Batch Normalization)	(None, 697)	2788
dropout_18 (Dropout)	(None, 697)	0
dense_58 (Dense)	(None, 458)	319684
batch_normalization_21 (Batch Normalization)	(None, 458)	1832
dropout_19 (Dropout)	(None, 458)	0
dense_59 (Dense)	(None, 246)	112914
batch_normalization_22 (Batch Normalization)	(None, 246)	984
dropout_20 (Dropout)	(None, 246)	0
dense_60 (Dense)	(None, 111)	27417
batch_normalization_23 (Batch Normalization)	(None, 111)	444
dropout_21 (Dropout)	(None, 111)	0
dense_61 (Dense)	(None, 58)	6496
batch_normalization_24 (Batch Normalization)	(None, 58)	232
dropout_22 (Dropout)	(None, 58)	0
dense_62 (Dense)	(None, 10)	590
Total params: 1,020,526		
Trainable params: 1,017,386		
Non-trainable params: 3,140		

```
In [0]: model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy'])
```

```
history = model.fit(X_train, Y_train, batch_size = batch_size, epochs = nb_epoch, verbose=0)
```


Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 22s 361us/step - loss: 1.7259 - acc: 0.4318 - va
Epoch 2/20
60000/60000 [=====] - 19s 315us/step - loss: 0.7668 - acc: 0.7482 - va
Epoch 3/20
60000/60000 [=====] - 18s 305us/step - loss: 0.5112 - acc: 0.8477 - va
Epoch 4/20
60000/60000 [=====] - 18s 304us/step - loss: 0.4055 - acc: 0.8850 - va
Epoch 5/20
60000/60000 [=====] - 18s 306us/step - loss: 0.3332 - acc: 0.9080 - va
Epoch 6/20
60000/60000 [=====] - 18s 306us/step - loss: 0.2908 - acc: 0.9211 - va
Epoch 7/20
60000/60000 [=====] - 18s 305us/step - loss: 0.2566 - acc: 0.9311 - va
Epoch 8/20
60000/60000 [=====] - 19s 314us/step - loss: 0.2320 - acc: 0.9391 - va
Epoch 9/20
60000/60000 [=====] - 18s 304us/step - loss: 0.2208 - acc: 0.9429 - va
Epoch 10/20
60000/60000 [=====] - 18s 305us/step - loss: 0.2035 - acc: 0.9466 - va
Epoch 11/20
60000/60000 [=====] - 18s 304us/step - loss: 0.1856 - acc: 0.9511 - va
Epoch 12/20
60000/60000 [=====] - 18s 305us/step - loss: 0.1777 - acc: 0.9539 - va
Epoch 13/20
60000/60000 [=====] - 18s 306us/step - loss: 0.1676 - acc: 0.9565 - va
Epoch 14/20
60000/60000 [=====] - 18s 306us/step - loss: 0.1569 - acc: 0.9589 - va
Epoch 15/20
60000/60000 [=====] - 18s 305us/step - loss: 0.1544 - acc: 0.9611 - va
Epoch 16/20
60000/60000 [=====] - 18s 308us/step - loss: 0.1485 - acc: 0.9624 - va
Epoch 17/20
60000/60000 [=====] - 18s 306us/step - loss: 0.1435 - acc: 0.9630 - va
Epoch 18/20
60000/60000 [=====] - 18s 306us/step - loss: 0.1323 - acc: 0.9654 - va
Epoch 19/20
60000/60000 [=====] - 18s 307us/step - loss: 0.1278 - acc: 0.9675 - va
Epoch 20/20
60000/60000 [=====] - 18s 306us/step - loss: 0.1241 - acc: 0.9685 - va
```

```
In [0]: score_relu_bn_drop5 = model.evaluate(X_test, Y_test, verbose = 0)
        print('Test score:', score_relu_bn_drop5[0])
        print('Test accuracy:', score_relu_bn_drop5[1])
```

```
fig, ax = plt.subplots(1, 1)
```

```

ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, v

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of e

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07801552013880574

Test accuracy: 0.9801



```

In [0]: w_after = model.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')


plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3')


plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='c')
plt.xlabel('Hidden Layer 4')


plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='w')
plt.xlabel('Hidden Layer 5')



plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning: remove_na is
  kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning: remove_na is
  violin_data = remove_na(group_data)

```



Observations

1. As we can see when apply bn + dropout each and every model performs well.

Conclusions 1. We perform various different types of model architecture and with different layers and also plotted train and validation error graph. We observe that when applied BN + Dropout with three different architecture models performs very well. And, when apply only dropout with 5 hidden layer, BN with 3 hidden layer, it works worse than all. 2. We also check for the weights of model and found that weights are slightly small in BN and Dropout, when applied individually. Other than this, model with only I/O layer also performs quite well. 3. We can also see the model performance chart in below table.

```
In [0]: from prettytable import PrettyTable
```

```
x = PrettyTable()
```

```
x.field_names = ["MODEL", "ACCURACY"]
```

```
x.add_row(["MLP with only I/O layer", score_in_out[1]])
x.add_row(["MLP with 2 hidden layer", score_relu2_adam[1]])
x.add_row(["MLP with 2 hidden layer", score_relu2_adam_diff[1]])
x.add_row(["MLP with 3 hidden layer", score_relu3_adam[1]])
x.add_row(["MLP + BN with 2 hidden layer", score_relu5_adam[1]])
x.add_row(["MLP + BN with 2 hidden layer", score_relu_bn2[1]])
x.add_row(["MLP + BN with 3 hidden layer", score_relu_bn2_diff[1]])
x.add_row(["MLP + BN with 5 hidden layer", score_relu_bn3[1]])
```

```

x.add_row(["MLP + Dropout with 2 hidden layer", score_relu_bn5[1]])
x.add_row(["MLP + Dropout with 2 hidden layer", score_relu_drop2[1]])
x.add_row(["MLP + Dropout with 3 hidden layer", score_relu_drop3[1]])
x.add_row(["MLP + Dropout with 5 hidden layer", score_relu_drop5[1]])
x.add_row(["MLP + BN + Dropout with 2 hidden layers", score_relu_bn_drop2[1]])
x.add_row(["MLP + BN + Dropout with 2 hidden layers", score_relu_bn_drop2_diff[1]])
x.add_row(["MLP + BN + Dropout with 3 hidden layers", score_relu_bn_drop3[1]])
x.add_row(["MLP + BN + Dropout with 5 hidden layers", score_relu_bn_drop5[1]])

print(x)

```

MODEL	ACCURACY
MLP with only I/O layer	0.9093
MLP with 2 hidden layer	0.9807
MLP with 2 hidden layer	0.9845
MLP with 3 hidden layer	0.9773
MLP + BN with 2 hidden layer	0.9787
MLP + BN with 2 hidden layer	0.9796
MLP + BN with 3 hidden layer	0.8002
MLP + BN with 5 hidden layer	0.9798
MLP + Dropout with 2 hidden layer	0.9815
MLP + Dropout with 2 hidden layer	0.9783
MLP + Dropout with 3 hidden layer	0.9328
MLP + Dropout with 5 hidden layer	0.0982
MLP + BN + Dropout with 2 hidden layers	0.9831
MLP + BN + Dropout with 2 hidden layers	0.982
MLP + BN + Dropout with 3 hidden layers	0.9824
MLP + BN + Dropout with 5 hidden layers	0.9801